

Probability in Machine Learning

Author

Marc'Antonio Lopez

AI & Data Analytics student at Polytechnic University of Turin

Probability theory forms the bedrock of machine learning, providing the essential mathematical framework for understanding and managing **uncertainty**. This enables well-informed decisions even with incomplete or unpredictable information. Moreover, it allows machine learning models to quantify belief, predict outcomes, and learn from inherently noisy or variable data.

Introduction to Probability and Random Events in Machine Learning

Why Study Probability?

Studying probability is crucial as it equips us with powerful methods for making predictions and decisions amidst **uncertainty**. Real-world phenomena are incredibly complex, influenced by countless often unknown or imprecisely measurable factors. This complexity renders purely **deterministic models** (where outcomes are always precisely predictable) impractical or impossible to build. Therefore, we utilize the concept of **random events** to quantify and systematically navigate this inherent complexity.

Types of Events

Understanding the nature of an event is fundamental:

- **Deterministic Event:** An event whose outcome is entirely predictable if all initial conditions and influencing factors are fully known. For example, the outcome of a simple mathematical calculation is deterministic.
- **Random Event:** An event whose specific outcome is uncertain and cannot be predicted with absolute certainty, even if some initial conditions are known. For example, the result of a coin flip or a die roll is a random event.

Different Ways to Interpret Probability

Probability can be understood and applied from several perspectives:

1. **Classical Interpretation:** This interpretation defines probability as the ratio of "favorable outcomes" to the "total number of equally likely possible outcomes."
 - **Example:** When rolling a fair six-sided die, the probability of rolling a 4 is $1/6$, because there is one favorable outcome (rolling a 4) out of six equally likely total outcomes.
2. **Frequentist Interpretation:** This view defines probability as the long-run relative frequency of an event. It's based on observing the outcome over a large number of trials.
 - **Example:** If you flip a fair coin many, many times, the proportion of times it lands on heads will eventually approach 0.5. The probability of heads is thus interpreted as 0.5.
3. **Bayesian Interpretation:** This is a more subjective interpretation, where probability represents a "degree of belief" in an event. This belief is updated as new evidence or data becomes available.
 - **Example:** Your initial belief about whether a new drug is effective might change (be updated) after seeing results from clinical trials.
4. **Axiomatic Treatment:** This is the formal mathematical approach to probability. It defines probability through a set of fundamental axioms (rules) that must be satisfied, allowing for a rigorous and consistent theory, regardless of the interpretation.

Example: Rolling a Standard Six-Sided Die

Let's use the example of rolling a fair, standard six-sided die to illustrate basic probability concepts:

1. **Sample Space (Ω):** This is the set of all possible outcomes of the experiment.
 - For a six-sided die, $\Omega = \{1, 2, 3, 4, 5, 6\}$.
2. **Individual Probabilities:** For a fair die, each individual outcome is equally likely.
 - The probability of rolling any specific number x is $P(\{x\}) = \frac{1}{6}$ for $x \in \{1, \dots, 6\}$.
3. **Compound Events:** To find the probability of an event that consists of multiple outcomes (a "compound event"), you sum the probabilities of its individual outcomes.
 - **Example:** Let's find the probability of rolling an "even number." The outcomes for this event are $\{2, 4, 6\}$.

$$P(\text{"even number"}) = P(\{2\}) + P(\{4\}) + P(\{6\}) = \frac{1}{6} + \frac{1}{6} + \frac{1}{6} = \frac{3}{6} = \frac{1}{2}$$

Probability Spaces and Axioms: The Formal

Framework

To provide a rigorous mathematical foundation for probability, we define a **probability space**, which is composed of three essential parts:

1. **Sample Space (Ω):** The complete set of all possible individual outcomes of a random experiment.
 - **Example:** For a single die roll, $\Omega = \{1, 2, 3, 4, 5, 6\}$.
2. **Event Space ($\mathcal{A}$ or $\mathcal{F}$):** A collection of specific subsets of the sample space (Ω). Each subset is an "event" for which we can calculate a probability. This collection must satisfy certain mathematical properties to form a σ -field (or σ -algebra), ensuring consistency.
3. **Probability Function (P):** A function assigning a numerical probability value (between 0 and 1, inclusive) to each event within the event space ($\mathcal{A}$).

Fundamental Properties (Axioms) of Probability

A valid probability function P must strictly adhere to **Kolmogorov's Axioms**, which are the cornerstone of probability theory:

1. **Non-negativity:** The probability of any event A_i must be non-negative (greater than or equal to zero).
 - $P(A_i) \geq 0$
2. **Normalization:** The probability of the entire sample space Ω (meaning something definitely happens) must be exactly 1.
 - $P(\Omega) = 1$
3. **Countable Additivity:** If you have a sequence of events $A_1, A_2, \dots$ that are **mutually exclusive** (meaning no two events can occur at the same time), then the probability that *any* of these events occur (their union) is equal to the sum of their individual probabilities.
 - $P\left(\bigcup_{n=1}^{\infty} A_n\right) = \sum_{n=1}^{\infty} P(A_n)$

Derived Properties of Probability

From Kolmogorov's Axioms, several other important properties of probability can be logically derived:

- **Probability of the Impossible Event:** The probability of an event that cannot occur (represented by the empty set $\emptyset$) is 0.
 - $P(\emptyset) = 0$
- **Probability of a Complement:** The probability of an event *not* occurring (A^C , the complement of event A) is 1 minus the probability of the event occurring.
 - $P(A^C) = 1 - P(A)$

- **Probability Range:** The probability of any event A must always fall within the range of 0 to 1, inclusive.
 - $0 \leq P(A) \leq 1$

Example: Rolling a Die (Revisited with Axioms)

Let's re-examine the die roll example through the lens of these axioms:

- **Axiom 1 (Non-negativity):** For a fair die, the probability of rolling any specific number x is $P(\{x\}) = 1/6$, which is indeed ≥ 0 . This axiom is satisfied.
- **Axiom 2 (Normalization):** The sum of the probabilities of all individual outcomes in the sample space is:

$$P(\Omega) = P(\{1\}) + P(\{2\}) + P(\{3\}) + P(\{4\}) + P(\{5\}) + P(\{6\}) = 1/6 + 1/6 + 1/6 + 1/6 + 1/6 + 1/6 = 6/6 = 1.$$
 This axiom is satisfied.
- **Axiom 3 (Countable Additivity):** Consider the event "even number," which comprises the outcomes $\{2, 4, 6\}$. Since the events $\{2\}$, $\{4\}$, and $\{6\}$ are mutually exclusive (you can't roll a 2 and a 4 simultaneously), we can use this axiom:

$$P(\text{"even number"}) = P(\{2\} \cup \{4\} \cup \{6\}) = P(\{2\}) + P(\{4\}) + P(\{6\}) = 1/6 + 1/6 + 1/6 = 3/6 = 1/2.$$
 This axiom is satisfied.

Conditional Probability and Independence

Conditional Probability: Probability Given New Information

Conditional probability, denoted as $P(A|B)$, is the probability of event A occurring *given that event B has already occurred*. This concept allows us to update our belief about an event's likelihood based on new information, and is defined only if $P(B) > 0$. The formula for conditional probability is given by:

$$P(A|B) = \frac{P(A \cap B)}{P(B)}$$

Here, $P(A \cap B)$ represents the probability that *both* event A and event B occur.

- **Example (Die Roll):**
 - Let **Event A** be rolling a 2 or a 3. So, $A = \{2, 3\}$. The probability of A is $P(A) = 2/6 = 1/3$.

- Let **Event B** be rolling a number greater than 1. So, $B = \{2, 3, 4, 5, 6\}$. The probability of B is $P(B) = 5/6$.
- The **intersection** of A and B ($A \cap B$) is the set of outcomes where both A and B occur: $A \cap B = \{2, 3\}$. The probability of this intersection is $P(A \cap B) = 2/6$.

Now, let's calculate the probability of Event A given Event B has occurred, $P(A|B)$:

$$P(A|B) = \frac{P(A \cap B)}{P(B)} = \frac{2/6}{5/6} = \frac{2}{5}$$

Interpretation: This result means that *if we already know* the die roll was greater than 1, the probability of that roll being a 2 or a 3 is now 2/5 (or 40%). This is higher than the original probability of A (1/3 or ~33.3%), indicating that knowing B occurred changes our assessment of A.

Bayes' Formula: Reversing the Condition


Bayes' Formula (also known as Bayes' Theorem) is one of the most fundamental and powerful results in probability theory, especially in machine learning. It allows us to "reverse" the conditionality; that is, to calculate $P(B|A)$ if we know $P(A|B)$, along with the individual probabilities $P(A)$ and $P(B)$.

The formula is:

$$P(B|A) = \frac{P(A|B)P(B)}{P(A)}$$

This theorem is central to Bayesian inference, where $P(B|A)$ is the "posterior probability" (updated belief), $P(A|B)$ is the "likelihood," $P(B)$ is the "prior probability," and $P(A)$ is the "evidence."

Here's a conceptual flow of how Bayes' Theorem works:



Independence: Events Not Influencing Each Other

Two events, A and B , are considered **statistically independent** if the occurrence of one event does not affect the probability of the other event occurring.

This relationship can be defined in several equivalent ways:

- **Primary Definition:** Events A and B are independent if the probability of both events occurring is equal to the product of their individual probabilities:
$$P(A \cap B) = P(A)P(B).$$
- **Equivalent Conditions (if $P(A) > 0$ and $P(B) > 0$):**
 - The probability of A given B is simply the probability of A : $P(A|B) = P(A)$.
 - The probability of B given A is simply the probability of B : $P(B|A) = P(B)$.

Verification using Bayes' Formula (Die Example)

Let's use the previous die roll example values to check for independence using Bayes' Formula. We had:

- $P(A|B) = 2/5$
- $P(A) = 1/3$
- $P(B) = 5/6$

Now, let's calculate $P(B|A)$ using Bayes' Formula:

$$P(B|A) = \frac{P(A|B)P(B)}{P(A)} = \frac{(2/5) \cdot (5/6)}{1/3} = \frac{10/30}{1/3} = \frac{1/3}{1/3} = 1$$

Interpretation of the Result: The result $P(B|A) = 1$ means that if we know the roll was a 2 or a 3 (Event A), it is *certain* (probability of 1) that the roll was also greater than 1 (Event B).

Conclusion on Independence: Since $P(B|A) = 1$ which is **not equal** to $P(B) = 5/6$, events A and B are **not independent**. This makes sense intuitively: knowing the roll was a 2 or 3 definitely tells us something about whether it was greater than 1.

Random Variables

A **random variable (RV)** is a powerful probability concept, serving as a bridge between abstract random experiment outcomes and numerical values. It maps each possible outcome to a real number, enabling mathematical analysis of uncertain situations.

Definition of a Random Variable

Formally, a random variable X is a **function** $X : \Omega \rightarrow \mathbb{R}$. This function assigns a unique real number, $X(\omega)$, to every individual outcome ω in the sample space Ω .

A crucial condition for a function to be considered a random variable is that for any real number x , the set of all outcomes ω for which $X(\omega) \leq x$ must constitute an event (i.e., $\{\omega : X(\omega) \leq x\}$ must be a member of the event space $\mathcal{F}$). This condition ensures the variable is "measurable," meaning we can meaningfully assign probabilities to its values.

Cumulative Distribution Function (CDF)

The **Cumulative Distribution Function (CDF)**, denoted as $F_X(x)$, is a fundamental function that completely describes the probability distribution of a random variable X . It gives the probability that the random variable X will take on a value less than or equal to a given number x .

The definition is: $F_X(x) = P(X \leq x)$.

- **Key Properties of a CDF:**

- i. **Bounded:** The probability value must always be between 0 and 1, inclusive.

- $0 \leq F_X(x) \leq 1$

- ii. **Non-decreasing:** As x increases, the CDF value must either stay the same or increase. It can never decrease.

- If $a < b$, then $F_X(a) \leq F_X(b)$.

iii. Limits at Extremes:

- As x approaches negative infinity, the CDF approaches 0 (no chance of being less than an infinitely small number).
 - $\lim_{x \rightarrow -\infty} F_X(x) = 0$
- As x approaches positive infinity, the CDF approaches 1 (certainty of being less than an infinitely large number).
 - $\lim_{x \rightarrow \infty} F_X(x) = 1$

iv. **Right-continuous:** The CDF must be continuous from the right. This means that as you approach a point x from values slightly larger than x , the function value approaches $F_X(x)$.

- $\lim_{h \rightarrow 0^+} F_X(x + h) = F_X(x)$

Types of Random Variables

Random variables are primarily categorized into two types based on the nature of the values they can take:

1. Discrete Random Variables

- **Definition:** These are random variables that can take on a **finite** number of distinct values, or a **countably infinite** number of distinct values (like integers: 0, 1, 2, ...). They typically represent counts or categories.
- **Description:** They are described by a **Probability Mass Function (PMF)**, denoted as $f_X(x)$. The PMF gives the probability that the random variable X is exactly equal to a specific value x .
 - $f_X(x) = P(X = x)$
- **Properties of a PMF:**
 - Each probability must be non-negative: $f_X(x) \geq 0$.
 - The sum of all probabilities for all possible values of x must equal 1: $\sum_{\text{all } x} f_X(x) = 1$.
- **Example:** The outcome of a fair die roll is a discrete random variable $X \in \{1, \dots, 6\}$. Its PMF is $f_X(x) = 1/6$ for each possible value x .

2. Continuous Random Variables

- **Definition:** These are random variables that can take on **any value within a continuous range** of real numbers. They typically represent measurements.
- **Description:** They are described by a **Probability Density Function (PDF)**, denoted as $f_X(x)$.
- **Key Concept:** For a continuous random variable, the probability of it taking on any *exact* single value is 0 ($P(X = x) = 0$). Instead, probability is calculated as the **area under the PDF curve**.

over a given interval.

- $P(a \leq X \leq b) = \int_a^b f_X(x)dx$

- **Properties of a PDF:**

- The function value must be non-negative: $f_X(x) \geq 0$. (Note: Unlike PMF, $f_X(x)$ itself can be greater than 1, as long as the integral over the entire range is 1).
- The total area under the entire PDF curve must equal 1: $\int_{-\infty}^{\infty} f_X(x)dx = 1$.

- **Relationship to CDF:**

- The CDF is the integral of the PDF: $F_X(x) = \int_{-\infty}^x f_X(t)dt$.
- Conversely, the PDF is the derivative of the CDF: $f_X(x) = \frac{d}{dx}F_X(x)$.

- **Example:** The Gaussian (Normal) Distribution is a common example of a continuous probability distribution, often used to model natural phenomena like height or measurement errors.

Random Vectors (Multidimensional Random Variables)

A **random vector** is a collection of multiple random variables, typically denoted as $X = (X_1, X_2, \dots, X_m)$. Each X_i is a random variable, and collectively they describe the outcomes of a multi-dimensional random experiment.

Joint Cumulative Distribution Function (Joint CDF)

The **Joint CDF** for a random vector $X = (X_1, \dots, X_m)$ describes the probability that *each component* X_i simultaneously takes on a value less than or equal to a corresponding specified value x_i .

The definition is:

$$F_X(x_1, \dots, x_m) = P(X_1 \leq x_1, \dots, X_m \leq x_m)$$

Marginal Distributions

Marginal distributions are derived from a joint distribution when you are interested in the probability distribution of a subset of the random variables, ignoring the values of the others. This is achieved by "integrating out" (for continuous variables) or "summing out" (for discrete variables) the unwanted variables.

- **For Marginal CDF $F_{X_i}(x_i)$:** You obtain the marginal CDF of a single variable X_i by taking the limit of the joint CDF as the values of all other variables x_j (where $j \neq i$) approach positive infinity. This effectively accounts for all possible values of the other variables.
 - $F_{X_i}(x_i) = \lim_{x_j \rightarrow \infty \text{ for all } j \neq i} F_X(x_1, \dots, x_m)$
- **Marginal PMFs/PDFs:** Similarly, marginal probability mass functions (PMFs) for discrete variables and marginal probability density functions (PDFs) for continuous variables are obtained by summing (for discrete) or integrating (for continuous) the joint PMFs/PDFs over all possible values of the variables you want to "ignore."

Statistical Independence of Random Variables

Two random variables, X and Y , are considered **statistically independent** if knowing the value of one variable provides absolutely no information about the value of the other variable.

This can be formally defined using their distributions:

- **Definition using CDFs:** X and Y are independent if their joint CDF is equal to the product of their individual (marginal) CDFs for all possible values of x and y .
 - $F_{X,Y}(x, y) = F_X(x)F_Y(y)$ for all x, y .
- **Definition using PDFs/PMFs:**
 - **For continuous variables:** Their joint PDF is equal to the product of their individual (marginal) PDFs.
 - $f_{X,Y}(x, y) = f_X(x)f_Y(y)$
 - **For discrete variables:** The probability of them taking on specific values x and y simultaneously is equal to the product of their individual (marginal) probabilities.
 - $P(X = x, Y = y) = P(X = x)P(Y = y)$

Conditional Probability Density Function (Conditional PDF)

The **Conditional PDF**, denoted as $f_{X|Y}(x|y)$, describes the probability distribution of random variable X *given that* random variable Y has taken on a specific value y . It is analogous to conditional probability for events but applied to continuous random variables.

The formula is:

$$f_{X|Y}(x|y) = \frac{f_{X,Y}(x, y)}{f_Y(y)} \quad (\text{where } f_Y(y) > 0)$$

Here, $f_{X,Y}(x, y)$ is the joint PDF of X and Y , and $f_Y(y)$ is the marginal PDF of Y . This formula

essentially normalizes the joint probability by the likelihood of the condition, giving us the probability density of X *within* that specific condition.

Expectations (Expected Values): Measures of Central Tendency and Spread

The **expectation** (or **expected value**) of a random variable, denoted as $\mathbb{E}[X]$ or μ_X , represents its long-run average value if the random experiment were repeated many times. Thus, it serves as a fundamental measure of **central tendency** for a probability distribution.

Definition of Expectation ($\mathbb{E}[X]$)

The calculation of expectation differs slightly for discrete and continuous random variables:

- **For a Discrete Random Variable (RV):** The expectation is the sum of each possible value multiplied by its probability mass function (PMF).
 - $\mathbb{E}[X] = \mu_X = \sum_{x \in S} x f_X(x)$
 - Where S is the set of all possible values for X .
- **For a Continuous Random Variable (RV):** The expectation is the integral of each possible value multiplied by its probability density function (PDF).
 - $\mathbb{E}[X] = \mu_X = \int_{-\infty}^{\infty} x f_X(x) dx$

Variance and Covariance: Measures of Spread and Relationship

While expectation tells us the central tendency, other measures help describe the spread or the relationship between variables.

1. Variance ($\text{Var}(X)$ or σ_X^2)

- **Purpose:** Variance is a measure of how widely dispersed the values of a random variable are around its mean. A high variance indicates that values tend to be spread out far from the mean, while a low variance indicates values are clustered closely around the mean.
- **Definition:** Variance is the expected value of the squared difference between the random variable and its mean.
 - $\text{Var}(X) = \sigma_X^2 = \mathbb{E}[(X - \mu_X)^2]$
- **Computational Formula:** A more convenient formula for calculation is:

- $\text{Var}(X) = \mathbb{E}[X^2] - (\mathbb{E}[X])^2$

- **Standard Deviation:** The standard deviation (σ_X) is simply the square root of the variance. It is often preferred because it has the same units as the random variable itself, making it easier to interpret.

- $\sigma_X = \sqrt{\text{Var}(X)}$

2. Covariance ($\text{Cov}(X, Y)$)

- **Purpose:** Covariance measures the extent to which two random variables, X and Y , change together.
 - **Positive covariance:** Indicates that X and Y tend to move in the same direction (e.g., as X increases, Y tends to increase).
 - **Negative covariance:** Indicates they tend to move in opposite directions (e.g., as X increases, Y tends to decrease).
 - **Near zero covariance:** Suggests little to no **linear** relationship between X and Y .
- **Definition:** Covariance is the expected value of the product of the deviations of X and Y from their respective means.
 - $\text{Cov}(X, Y) = \mathbb{E}[(X - \mu_X)(Y - \mu_Y)]$
- **Computational Formula:** A more convenient formula for calculation is:
 - $\text{Cov}(X, Y) = \mathbb{E}[XY] - \mathbb{E}[X]\mathbb{E}[Y]$
- **Important Note:** If two random variables X and Y are statistically independent, then their covariance will be 0. However, the reverse is not always true: a covariance of 0 only implies no *linear* relationship; it does not necessarily mean the variables are independent (they could have a non-linear relationship).

3. Correlation Coefficient ($\rho(X, Y)$)

- **Purpose:** The correlation coefficient is a normalized measure of the strength and direction of the **linear** relationship between two random variables X and Y . It scales the covariance to a standard range, making it easier to interpret.
- **Definition:**
 - $\rho(X, Y) = \frac{\text{Cov}(X, Y)}{\sigma_X \sigma_Y}$
- **Interpretation:** The correlation coefficient always ranges from -1 to +1:
 - **+1:** Signifies a perfect positive linear relationship (as one variable increases, the other increases proportionally).
 - **-1:** Indicates a perfect negative linear relationship (as one variable increases, the other decreases proportionally).
 - **0:** Implies no linear correlation between the variables.

Bayesian Estimation: Inferring Parameters with Prior Beliefs

Estimation is the process of using observed data to infer the unknown parameters of an underlying probability distribution or process. **Bayesian estimation** is a specific approach that uniquely incorporates prior beliefs or knowledge about these parameters into the inference process.

Maximum Likelihood Estimation (MLE) - A Frequentist Perspective

Before delving into Bayesian estimation, it's helpful to understand **Maximum Likelihood Estimation (MLE)**, which is a common frequentist approach.

- **Principle:** MLE identifies the values for the unknown parameters that make the observed data *most probable* (i.e., maximize the likelihood of the data given the parameters). It focuses solely on the likelihood of the observed data.
- **Example (Bernoulli Distribution):** Imagine you flip a coin n times and observe x_i (where $x_i = 1$ for heads, $x_i = 0$ for tails). You want to estimate the probability of success (π) for this coin.
 - The MLE estimate for π (π^{ML}) is simply the sample proportion of successes:

$$\pi^{\text{ML}} = \frac{\text{number of successes}}{\text{number of trials}} = \frac{1}{n} \sum_{i=1}^n x_i$$

- For instance, if you flip a coin 10 times and get 7 heads, the MLE for π would be $7/10 = 0.7$.

The Bayesian Approach: Incorporating Prior Beliefs

In contrast to MLE, the Bayesian approach treats the unknown parameter itself (e.g., π for the coin) as a **random variable**. This means we assign a probability distribution to the parameter to reflect our uncertainty about its true value.

The Bayesian estimation process involves combining three key components:

1. **Prior Distribution ($f_{\Pi}(\pi)$):** This distribution represents your initial beliefs or knowledge about the parameter π *before* you observe any data. It's your "best guess" or assumption about the parameter's possible values and their likelihoods.
2. **Likelihood ($P(X = x|\pi)$):** This is the probability of observing your specific dataset x given a hypothetical (or assumed) value for the parameter π . This is the same likelihood function used in

MLE.

3. **Posterior Distribution ($f_{\Pi|X}(\pi|x)$):** This is the central outcome of Bayesian estimation. It represents your **updated beliefs** about the parameter π *after* observing the data. It combines your initial prior beliefs with the evidence provided by the observed data through **Bayes' Theorem**:

$$f_{\Pi|X}(\pi|x) = \frac{P(X = x|\pi)f_{\Pi}(\pi)}{\int P(X = x|\pi')f_{\Pi}(\pi')d\pi'}$$

The denominator, $\int P(X = x|\pi')f_{\Pi}(\pi')d\pi'$, is a normalizing constant (often called "evidence" or "marginal likelihood") that ensures the posterior distribution integrates to 1. The posterior distribution provides a complete probabilistic summary of what you now believe about the parameter, given both your prior knowledge and the new data.

Density Estimation: Uncovering Data's Underlying Structure

Density estimation is the process of building a model to estimate the underlying probability density function (PDF) or probability mass function (PMF) from a given set of observed data. Essentially, it tries to understand the true distribution from which the data was sampled.

The Gaussian (Normal) Distribution: A Ubiquitous Example

The **Gaussian distribution** (also known as the **Normal distribution**) is a continuous probability distribution that is exceptionally important and frequently encountered in statistics and machine learning. Its prominence stems from its useful mathematical properties and its natural occurrence in many real-world phenomena.

- **Probability Density Function (PDF):** The mathematical formula that defines the shape of the Gaussian distribution is:

$$f_X(x; \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

Here:

- μ (mu) is the **mean** of the distribution, which represents its center.

- σ^2 (sigma squared) is the **variance** of the distribution, which represents its spread.
- π is the mathematical constant (approximately 3.14159).
- e is Euler's number (the base of the natural logarithm, approximately 2.71828).
- **Central Limit Theorem (CLT):** One of the most significant properties explaining the Gaussian distribution's ubiquity is the **Central Limit Theorem**. This theorem states that the sum or average of a large number of independent and identically distributed random variables will tend to follow a Gaussian distribution, *regardless of the original distribution of the individual variables*. This is why Gaussian distributions often appear in aggregate measurements.

Maximum Likelihood Estimation (MLE) for Gaussian Parameters

When you have observed data points $(x_1, \dots, x_n)$ that are assumed to be independent and identically distributed (i.i.d.) samples from an underlying Gaussian distribution $N(\mu, \sigma^2)$, you can use Maximum Likelihood Estimation (MLE) to estimate its parameters (μ and σ^2).

The MLE estimates for the mean (μ) and the precision ($\lambda = 1/\sigma^2$) are:

- **Mean (μ^{ML}):** The MLE estimate for the mean is simply the **sample mean** of your observed data:

$$\mu^{\text{ML}} = \bar{x} = \frac{1}{n} \sum_{i=1}^n x_i$$

- **Precision (λ^{ML}):** The MLE estimate for the precision (the reciprocal of variance) is:

$$\lambda^{\text{ML}} = \frac{n}{\sum_{i=1}^n (x_i - \mu^{\text{ML}})^2}$$

- **Variance ($\sigma^{2\text{ML}}$):** From the precision, the MLE estimate for the variance is:

$$\sigma^{2\text{ML}} = \frac{1}{\lambda^{\text{ML}}} = \frac{1}{n} \sum_{i=1}^n (x_i - \mu^{\text{ML}})^2$$

(Note: This MLE estimate for the variance is known to be **slightly biased** for small sample sizes. A common unbiased estimator for variance uses $n-1$ in the denominator instead of n .)

Visual Enhancements for Conceptual Understanding

The following diagrams and tables are designed to help visualize and summarize the key probability concepts discussed.

Conceptual Flow Diagram of Probability Concepts

This Mermaid diagram illustrates the logical progression and interconnectedness of the various probability concepts covered:




Table: Common Probability Distributions

This table summarizes key characteristics of some common probability distributions, useful in machine learning:

Distribution	Parameter(s)	Type	PMF / PDF Expression	Typical Use Case
Bernoulli	p (probability of success)	Discrete	$P(X = x) = p^x(1 - p)^{1-x}$ for $x \in \{0, 1\}$	Modeling a single trial with exactly two possible outcomes (e.g., a coin flip, success/failure).
Binomial	n (number of trials), p (success probability)	Discrete	$P(X = x) = \binom{n}{x} p^x (1 - p)^{n-x}$ for $x \in \{0, \dots, n\}$	Modeling the number of successes in a fixed number (n) of independent Bernoulli trials.
Gaussian	μ (mean), σ^2 (variance)	Continuous	$f_X(x) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$	Widely used for modeling many natural phenomena (e.g., height, errors), and the sum/average of many random variables (due to the Central Limit Theorem).