

Robot Learning

Lecture 8 Bandits and the Exploration-Exploitation dilemma

News

- CPD
- Lecture on Dec 23
- Project assistance in Jan? (1.5h+?)

Outline

- Introduction
- Exploration and Exploitation
 - Simple naïve exploration (ϵ -greedy)
 - Optimistic approaches
 - Probability matching & Information Value
- Bandits
 - Multi-armed
 - Contextual
- Back to MDPs

Introduction

Exploration-Exploitation Dilemma

- Online decision-making involves a fundamental choice:
 - **Exploitation** - Make the best decision given current information
 - **Exploration** - Gather more information
- The best long-term strategy may involve short-term sacrifices
- Gather enough information to make the best overall decisions

Examples

- Restaurant Selection
 - **Exploitation** - Go to your favourite restaurant
 - **Exploration** - Try a new restaurant
- Holiday planning
 - **Exploitation** – The camping site you go to since you are born
 - **Exploration** – Hitchhike and follow the flow
- Game Playing
 - **Exploitation** - Play the move you believe is best
 - **Exploration** - Play an experimental move


Exploration/Exploitation visualized




Exploration/Exploitation visualized




Exploration/Exploitation visualized



Exploration/Exploitation visualized

Diamond present where your friend was digging	Diamond present where you were digging	Present somewhere else
		

Exploration/Exploitation visualized




Principles

- Random Exploration
 - Add noise to greedy policy (e.g. ϵ -greedy)
- Optimism in the Face of Uncertainty
 - Estimate uncertainty on value
 - Prefer to explore states/actions with highest uncertainty
- Information State Search
 - Consider agent's information (history) as part of its state
 - Lookahead to see how information helps rewards

Bandits

Multi-Armed Bandit

- A **multi-armed bandit** is a tuple $\langle \mathcal{A}, \mathcal{R} \rangle$
 - $\mathcal{A}$ is a **known set of m actions** (or “arms”)
 - $\mathcal{R}^a(r) = P(r|a)$ is an **unknown probability distribution over rewards**
- At each step t the agent selects an action $a \in \mathcal{A}$
- The environment generates a reward $r_t \sim \mathcal{R}^{a_t}$
- The goal is to maximise cumulative reward $\sum_{\tau=1}^t r_{\tau}$



Regret

- The **action-value** is the mean reward for action a

$$Q(a) = \mathbb{E}[r|a]$$

Instead of $Q(s, a)$, because we have a single state

- The **optimal value** V^* is

$$V^* = Q(a^*) = \max_{a \in \mathcal{A}} Q(a)$$

- The **regret** is the opportunity loss for one step

$$I_t = \mathbb{E}[V^* - Q(a_t)]$$

- The total regret is the total opportunity loss

$$L_t = \mathbb{E} \left[\sum_{\tau=1}^t V^* - Q(a_\tau) \right]$$

- Maximise cumulative reward $\equiv$ minimise total regret


Counting Regret

- The **count** $N_t(a)$ is the **expected number of selections for action a** (how many times we pull that specific lever)
- The **gap** (for action a) $\Delta_a = V^* - Q(a)$ is the **difference in value between action a and optimal action a^*** (zero for optimal action)
- Regret is a **function of gaps and the counts**

$$L_t = \mathbb{E} \left[\sum_{\tau=1}^t V^* - Q(a_\tau) \right] = \sum_{a \in \mathcal{A}} \mathbb{E}[N_t(a)](V^* - Q(a)) = \sum_{a \in \mathcal{A}} \mathbb{E}[N_t(a)]\Delta_a$$

- A good algorithm ensures **small counts for large gaps**
- **Problem:** gaps are not known! – because I don't know the optimal V^*

Linear or Sublinear Regret



- If an algorithm **forever explores** it will have linear total regret
- If an algorithm **never explores** it will have linear total regret
- Is it possible to **achieve sublinear total regret**?

Exploration Strategies

Greedy Algorithms

- We consider **algorithms that estimate** $\hat{Q}_t(a) \approx Q(a)$
- Estimate the value of each action by Monte-Carlo evaluation

$$\hat{Q}_t(a) = \frac{1}{N_t(a)} \sum_{\tau} r_{\tau} \mathbf{1}(a_{\tau}; a)$$

- The **greedy algorithm** selects action with highest value

$$a_t^* = \arg \max_{a \in \mathcal{A}} \hat{Q}_t(a)$$

- Greedy can lock onto a suboptimal action forever (could be mitigated with a very optimistic optimization)

Greedy has linear total regret

ϵ -greedy Algorithms

- The ϵ -greedy algorithm continues to explore forever
 - With probability $1 - \epsilon$ select $a = \max_{a \in \mathcal{A}} \hat{Q}(a)$
 - With probability ϵ select a random action
- Constant ϵ ensures minimum total regret L_t (lower bound)

$$L_t \geq \frac{\epsilon}{|\mathcal{A}|} \sum_{a \in \mathcal{A}} \Delta_a$$

Action space
size

ϵ -greedy has linear total regret

Decaying ϵ_t -greedy Algorithms

- Pick a decay schedule for $\epsilon_1, \epsilon_2, \dots$
- Consider the following schedule

$$c > 0$$

$$d = \min_{a | \Delta_a > 0} \Delta_i$$

$$\epsilon_t = \min \left\{ 1, \frac{c|\mathcal{A}|}{d^2 t} \right\}$$

Ensures
convergence

- Decaying ϵ_t -greedy has **logarithmic asymptotic total regret**
- Unfortunately, schedule requires **advance knowledge of gaps**
- **Goal**: find an algorithm with sublinear regret for any multi-armed bandit (without knowledge of R)

Lower Bound

- The performance of any algorithm is **determined by similarity between optimal arm and other arms**
- Hard problems have similar-looking arms with different means
- This is described formally by the gap Δ_a and the similarity in distributions $KL(\mathcal{R}^a || \mathcal{R}^{a^*})$

Lower Bound

- The performance of any algorithm is **determined by similarity between optimal arm and other arms**
- Hard problems have similar-looking arms with different means
- This is described formally by the gap Δ_a and the similarity in distributions $KL(\mathcal{R}^a || \mathcal{R}^{a^*})$


Theorem (Lai and Robbins)

Asymptotic total regret is at least logarithmic in the number of steps

$$\lim_{t \rightarrow \infty} L_t \geq \log t \sum_{a | \Delta_a > 0} \frac{\Delta_a}{KL(\mathcal{R}^a || \mathcal{R}^{a^*})}$$


Optimism in the Face of Uncertainty (I)

- Which action should we pick?
- The more uncertain we are about an action-value, the more important it is to explore that action
- It could turn out to be the best action



Optimism in the Face of Uncertainty (II)

- After picking blue action
- We are less uncertain about the value
- And more likely to pick another action
- Until we home in on best action



Upper Confidence Bounds

- Estimate an **upper confidence** $\hat{U}_t(a)$ for each action value
- Such that $Q(a) \leq \hat{Q}_t(a) + \hat{U}_t(a)$ with high probability
- This depends on the number of times $N(a)$ has been selected
 - **Small** $N_t(a) \Rightarrow$ large $\hat{U}_t(a)$ (estimated value is **uncertain**)
 - **Large** $N_t(a) \Rightarrow$ small $\hat{U}_t(a)$ (estimated value is **accurate**)
- Select action maximising **Upper Confidence Bound** (UCB)

$$a_t = \arg \max_{a \in \mathcal{A}} \hat{Q}_t(a) + \hat{U}_t(a)$$

Hoeffding's Inequality

Theorem (Hoeffding's Inequality)

Let $X_1, \dots, X_t$ be i.i.d. random variables in $[0,1]$ and let $\bar{X}_t = \frac{1}{t} \sum_{\tau=1}^t X_\tau$ the sample mean, then

$$P(\mathbb{E}[X] > \bar{X}_t + u) \leq e^{-2tu^2}$$

We will apply Hoeffding's Inequality to rewards of the bandit conditioned on selecting action a

$$P(Q(a) > \hat{Q}_t(a) + U_t(a)) \leq e^{-2N_t(a)U_t(a)^2}$$

Calculating Upper Confidence Bounds

- Pick a probability p that true value exceeds UCB
- Now solve for $U_t(a)$

$$e^{-2N_t(a)U_t(a)^2} = p$$
$$U_t(a) = \sqrt{\frac{-\log p}{2N_t(a)}}$$

- Reduce p as we observe more rewards, e.g. $p = t^{-4}$
- Ensures we select optimal action as $t \rightarrow \infty$

$$U_t(a) = \sqrt{\frac{2\log t}{N_t(a)}}$$

UCB1

- This leads to the **UCB1 algorithm**

$$a_t = \arg \max_{a \in \mathcal{A}} Q(a) + \sqrt{\frac{2 \log t}{N_t(a)}}$$

- The UCB algorithm achieves **logarithmic asymptotic total regret**

UCB vs. ϵ - greedy on 10- armed Bandit





Figure 9. Comparison on distribution 11 (10 machines with parameters 0.9, 0.6, ..., 0.6).



Bayesian Bandits

- So far, no assumptions about the reward distribution $\mathcal{R}$
 - Except bounds on rewards
- Bayesian bandits exploit prior knowledge of rewards $P(\mathcal{R})$
- They compute posterior distribution of rewards $P(\mathcal{R}|h_t)$
 - where $h_t = a_1, r_1; \dots; a_{t-1} r_{t-1}$ is the history
- Use posterior to guide exploration
 - Upper confidence bounds (Bayesian UCB)
 - Probability matching (Thompson sampling)
- Better performance if prior knowledge is accurate


Bayesian UCB Example - Independent Gaussians

Assume reward distribution is Gaussian

$$\mathcal{R}_a(r) = \mathcal{N}(r; \mu_a, \sigma_a^2)$$

- Compute Gaussian posterior over μ_a and σ_a^2 (Bayes)

$$P(\mu_a, \sigma_a^2 | h_t) \propto P(\mu_a, \sigma_a^2) \prod_{t|a_t=a} \mathcal{N}(r_t; \mu_a, \sigma_a^2)$$



Bayesian UCB Example - Independent Gaussians

Assume reward distribution is Gaussian


$$\mathcal{R}_a(r) = \mathcal{N}(r; \mu_a, \sigma_a^2)$$

- Compute Gaussian posterior over μ_a and σ_a^2 (Bayes)

Prior distribution

$$P(\mu_a, \sigma_a^2 | h_t) \propto P(\mu_a, \sigma_a^2) \prod_{t|a_t=a} \mathcal{N}(r_t; \mu_a, \sigma_a^2)$$

Likelihood of experiencing a reward r_t given the distribution



Bayesian UCB Example - Independent Gaussians

Assume reward distribution is Gaussian


$$\mathcal{R}_a(r) = \mathcal{N}(r; \mu_a, \sigma_a^2)$$

- Compute Gaussian posterior over μ_a and σ_a^2 (Bayes)

$$P(\mu_a, \sigma_a^2 | h_t) \propto P(\mu_a, \sigma_a^2) \prod_{t|a_t=a} \mathcal{N}(r_t; \mu_a, \sigma_a^2)$$

- Pick action that maximises standard deviation of $Q(a)$

$$a_t = \arg \max \mu_a + c \sigma_a / \sqrt{N(a)}$$



Probability Matching

- Probability matching selects **action a according to probability that a is the optimal action**

$$\pi(a|h_t) = P\left(Q(a) = \max_{a'} Q(a') | h_t\right)$$

- Probability matching is **optimistic** in the face of uncertainty
 - Uncertain actions have higher probability of being max
- Can be **difficult to compute analytically** from posterior
 - Needs to be approximated through sampling

Thompson Sampling

- Thompson sampling implements probability matching

$$\pi(a|h_t) = \mathbb{E}_{\mathcal{R}|h_t} \left[\mathbf{1} \left(Q(a); \arg \max_{a' \in \mathcal{A}} Q(a') \right) | h_t \right]$$

- Use **Bayes law to compute posterior** distribution $P(\mathcal{R}|h_t)$
- Sample a reward** distribution $\mathcal{R}$ from posterior
- Compute action-value function $Q(a) = \mathbb{E}[\mathcal{R}_a]$
- Select **action maximising value on sample** $a_t = \arg \max_{a' \in \mathcal{A}} Q(a)$
- Thompson sampling achieves Lai and Robbins lower bound!

Information State

Value of Information

- **Exploration** is useful because it **gains information**
- Can we **quantify the value of information**?
 - How much reward a decision-maker would be prepared to pay in order to have that information, prior to making a decision
 - Long-term reward after getting information - immediate reward
- Information **gain is higher in uncertain** situations
- Therefore, it makes sense to explore uncertain situations more
- If we know value of information, we can trade-off exploration and exploitation optimally

Information State Space

- We have viewed bandits as one-step decision-making problems
- Can also **view as sequential decision-making** problems
- At each step there is an information state $\tilde{s}$
 - $\tilde{s}$ is a statistic of the history, i.e. $\tilde{s} = f(h_t)$
 - summarizes all information accumulated so far
- Each **action a causes a transition to a new information state $\tilde{s}'$** (and adds information) with probability $\tilde{P}_{\tilde{s}, \tilde{s}'}^a$
- Defines an **MDP $\tilde{\mathcal{M}}$ in augmented information state space**

$$\tilde{\mathcal{M}} = \langle \tilde{\mathcal{S}}, \mathcal{A}, \tilde{\mathcal{P}}, \mathcal{R}, \gamma \rangle$$

Example - Bernoulli Bandits


- Consider a Bernoulli bandit, such that $R_a = B(\mu_a)$ (e.g. win or lose a game with probability μ_a)
- Want to find which arm has the highest μ_a
- The information state is $\tilde{s} = \langle \alpha, \beta \rangle$
 - α_a counts the pulls of arm a where reward was 0
 - β_a counts the pulls of arm a where reward was 1

Solving Information State Space Bandits

- We now have an infinite MDP over information states that can be solved by reinforcement learning
- **Model-free** reinforcement learning
 - e.g. Q-learning (Duff, 1994)
- **Bayesian model-based** reinforcement learning
 - e.g. Gittins indices (Gittins, 1979)
 - This approach is known as **Bayes-adaptive RL**
 - Finds **Bayes-optimal exploration/exploitation trade-off** with respect to prior distribution

Bayes-Adaptive Bernoulli Bandits

- Start with $Beta(\alpha_a, \beta_a)$ prior over reward function R_a
- Each time a is selected, update posterior for R_a
 - $Beta(\alpha_a + 1, \beta_a)$ if $r = 0$
 - $Beta(\alpha_a, \beta_a + 1)$ if $r = 1$
- This defines transition function $\tilde{P}$ for the Bayes-adaptive MDP
- Information state $\langle \alpha, \beta \rangle$ corresponds to reward model $Beta(\alpha, \beta)$
- Each state transition corresponds to a Bayesian model update



Gittins Indices for Bernoulli Bandits

- Bayes-adaptive MDP can be solved by dynamic programming
- The solution is known as the Gittins index
- Exact solution to Bayes-adaptive MDP is typically intractable
 - Information state space is too large
- More recent idea: apply simulation-based search (Guez et al. 2012)
 - Forward search in information state space
 - Using simulations from current information state

Contextual Bandits

Contextual Bandits

- A contextual bandit is a tuple $\langle \mathcal{A}, \mathcal{S}, \mathcal{R} \rangle$
- $\mathcal{S} = P(s)$ is an unknown distribution over states (**contexts**)
- $\mathcal{R}_s^a(r) = P(r|s, a)$ is an **unknown probability distribution over rewards**
- At each step t
 - Environment generates state $s_t \sim \mathcal{S}$
 - Agent selects action $a_t \sim \mathcal{A}$
 - Environment generates reward $r_t \sim \mathcal{R}_{s_t}^{a_t}$



Linear Regression

- Action-value function is expected reward for state s and action a

$$Q(s, a) = \mathbb{E}[r|s, a]$$

- Estimate value function with a linear function approximator

$$Q_{\theta}(s, a) = \phi(s, a)^T \theta \approx Q(s, a)$$

- Estimate parameters by least squares regression

$$A_t = \sum_{\tau=1}^t \phi(s_{\tau}, a_{\tau}) \phi(s_{\tau}, a_{\tau})^T$$
$$b_t = \sum_{\tau=1}^t \phi(s_{\tau}, a_{\tau}) r_{\tau}$$
$$\theta_t = A_t^{-1} b_t$$


Linear Upper Confidence Bounds

- Least squares regression estimates the mean action-value $Q_{\theta}(s, a)$
- But it **can also estimate the variance** of the action-value $\sigma_{\theta}^2 = (s, a)$
 - i.e. the uncertainty due to parameter estimation error
- Add on a **bonus for uncertainty**, $U_{\theta}(s, a) = c\sigma$
 - i.e. define UCB to be c standard deviations above the mean

Geometric Interpretation

- Define **confidence ellipsoid** ε_t around parameters θ_t
- Such that ε_t **includes true parameters** θ^* with **high probability**
- Use this ellipsoid to estimate the uncertainty of action values
- Pick **parameters within ellipsoid that maximise action value**

$$\arg \max_{\theta \in \varepsilon} Q_{\theta}(s, a)$$



Calculating Linear Upper Confidence Bounds (LinUCB)

- For least squares regression, parameter covariance is A^{-1}
- Action-value is linear in features

$$Q_{\theta}(s, a) = \phi(s, a)^T \theta$$

- So, action-value **variance is quadratic**

$$\sigma_{\theta}^2(s, a) = \phi(s, a)^T A^{-1} \phi(s, a)$$

- Upper **confidence bound** is

$$Q_{\theta}(s, a) + c\sqrt{\phi(s, a)^T A^{-1} \phi(s, a)}$$

- Select **action maximizing upper confidence bound**

$$a_t = \arg \max_{a \in \mathcal{A}} Q_{\theta}(s_t, a) + c\sqrt{\phi(s, a)^T A^{-1} \phi(s, a)}$$

Exploration-Exploitation in MDPs

Applying Exploration/Exploitation to MDPs

The same principles for exploration/exploitation apply to MDPs

- Naive Exploration
- Optimism in the Face of Uncertainty
- Probability Matching
- Information State Search

Upper Confidence Bounds - Model-Free RL (Policy iteration)

- Maximise **UCB on action-value function** $Q^\pi(s, a)$
$$a_t = \arg \max_{a \in \mathcal{A}} Q(s_t, a) + U(s_t, a)$$
 - Estimate uncertainty in policy evaluation
 - Ignore uncertainty from policy improvement
- Maximise **UCB on optimal action-value function** $Q^*(s, a)$
$$a_t = \arg \max_{a \in \mathcal{A}} Q(s_t, a) + U_1(s_t, a) + U_2(s_t, a)$$
 - Estimate **uncertainty in policy evaluation**
 - Estimate **uncertainty from policy improvement**

Information State Search in MDPs

- MDPs can be **augmented to include information state**
- Now the augmented state is $\langle s, \tilde{s} \rangle$
 - s is original state within MDP
 - $\tilde{s}$ is a statistic of the history
- Each **action a causes a transition**
 - to a new state s' with probability $\mathcal{P}_{s,s'}^a$
 - to a new information state $\tilde{s}'$
- Defines MDP $\tilde{\mathcal{M}}$ in augmented information state space
$$\tilde{\mathcal{M}} = \langle \tilde{\mathcal{S}}, \mathcal{A}, \tilde{\mathcal{P}}, \mathcal{R}, \gamma \rangle$$
- **Posterior distribution** over MDP model is an **information state**
$$\tilde{s} = P(\mathcal{P}, \mathcal{R} | h_t)$$
- Solve this MDP to find **optimal exploration/exploitation trade-off** (w.r.t. prior)

Wrap-up

Take home messages

- A selection of principles for exploration/exploitation
 - Naive methods (ϵ -greedy)
 - Upper confidence bounds
 - Probability matching
 - Information state search
- Principles developed in bandit setting but also apply to MDP setting

Coming up

- Sim-to-real transfer
 - How can I manage to learn in simulation and transfer a policy in the real world
- Imitation Learning
- Research Topics
- Thesis ideas