

# Robot Learning

## Lecture 10: Supervised RL (Imitation, Inverse)

# Lecture 10:

Supervised robot learning: Behavior Cloning, Imitation, Generative IL, Inverse RL

Slides partially adapted from Pieter Abbeel (Stanford) Davide Bacciu (Unipi)

# Outline

- Introduction
- Imitation learning challenges
  - Distribution mismatch
  - Sequential models
  - Multimodal actions
- Advanced topics
  - Generative imitation learning
  - Inverse reinforcement learning

# Introduction

# Limitations of learning by (physical) interaction



The agent should have the chance to try (and fail) MANY times

- ✓ Hard when safety is a concern
- ✓ Hard in general when each interaction takes time

# Limitations of learning by (physical) interaction



# Imitation Learning

## Learning from demonstrations





- ✓ Kinesthetic imitation
  - ✓ Teacher takes over the end effectors of the agent.
  - ✓ Demonstrated actions are in the action space of the imitator



- ✓ Visual imitation
  - ✓ The actions of the teacher need to be inferred from visual sensory input and mapped to the action space of the agent



# Imitation Learning Vs Supervised Labelling







# Imitation Learning Vs Supervised Labelling



Our actions influence  
future state and data




Predicted labels do  
not influence future




# Action Labelling

A mapping from states/observations to action labels




Assume action labels in an annotated video are i.i.d

Train a classifier to map observations to labels at each time step




# Imitation Learning (Behaviour Cloning)

Policy - A mapping from observations to actions





Assume **actions in the expert** trajectories are i.i.d

Train a function to map observations to actions at each time step



# What Possibly Can Go Wrong?

- ✓ **Compounding** errors
  - ✓ Data augmentation
- ✓ **Non-markovian** observations
  - ✓ Recurrent models
- ✓ **Stochastic** expert actions
  - ✓ Generative modelling



# Distribution Shifts



# Independent in Time Error

At each time step  $t$ , the agent **wakes up on a state drawn from the state distribution of the expert trajectories**, and executes an action

- ✓ Error at each time  $t$  step bounded by  $\epsilon$
- ✓ Expected total error for  $T$  steps:  
 $\mathbb{E}[E] \leq \epsilon T$





# Compounding Errors

At each time step  $t$ , the agent **wakes up** on a state drawn from the state distribution resulting from executing the action suggested by the learned policy previously

- ✓ Error at each time  $t$  step bounded by  $\epsilon$
- ✓ Expected total error for  $T$  steps:


$$\mathbb{E}[E] \leq \epsilon(T + (T - 1) + (T - 2) + \dots) \propto \epsilon T^2$$



# Distribution Mismatch (Distribution Shift)


Due to the interdependence between our action at time step  $t$  and the state at  $t + 1$ , **states seen at test time** may come from a **different distribution than those seen at training time**

$$P_{\pi^*}(S_t) \neq P_{\pi_\theta}(S_t)$$



# Something Similar Happens in NLP

Sequence  
generation/transduction



Solutions use **teacher forcing**  
with **annealed schedule**

# Scheduled Sampling

- ✓ Initial training phase
  - ✓ Conditioning states come from the Teacher (training data)
- ✓ Later training phase
  - ✓ States/observations are sampled from the output of the model
- ✓ Ground-truth for the next time step still from the expert
  - ✓ Model learns to handle its mistakes
  - ✓ Pushing deviating generated sequences back to the right track

# Distribution Mismatch (DM)

|       | Supervised Learning | Supervised learning<br>+ Control      |
|-------|---------------------|---------------------------------------|
| Train | $(x, y) \sim P(D)$  | $s_t \sim P_{\pi^*}(\mathbf{s})$      |
| Test  | $(x, y) \sim P(D)$  | $s_t \sim P_{\pi_\theta}(\mathbf{s})$ |

Fundamental assumption in (standard) supervised learning is that training and test data distributions match




# DM Solution - Augment Data by querying the expert

- ✓ Change  $P_{\pi^*}(\mathbf{s})$  by augmenting the expert demonstration trajectories
  - ✓ Add examples in expert demonstration trajectories to cover the states where the agent will land when trying out its own policy
- ✓ Approaches
  - ✓ Generate synthetic data in simulation (ALVINN 1989)
  - ✓ Collecting additional data via clever tricks
  - ✓ Interactively query the experts in additional datapoints

# Clever Tricks – NVIDIA 2016

*End to End Learning for Self-Driving Cars , Bojarski et al. 2016*



Additional, left and right cameras  
with automatic ground-truth labels  
to recover from mistakes

# NVIDIA 2016 Demo

---

## DAVE 2 Driving a Lincoln

- A convolutional neural network
- Trained by human drivers
- Learns perception, path planning, and control  
"pixel in, action out"
- Front-facing camera is the only sensor

<https://youtu.be/NJU9ULQUwng>

# Incremental Dataset Growing - DAgger

How can we make  $P_{\pi^*}(\mathbf{s}) \approx P_{\pi_\theta}(\mathbf{s})$ ?

**Key Idea** - If we cannot be clever on  $P_{\pi_\theta}(\mathbf{s})$  let us be clever on  $P_{\pi^*}(\mathbf{s})$

## DAgger – Dataset Aggregation

- ✓ Collect training data from  $P_{\pi_\theta}(\mathbf{s})$  in place of  $P_{\pi^*}(\mathbf{s})$
- ✓ Collect observations by running  $\pi_\theta(a_t|s_t)$  and ask someone for labels  $a_t$

# Incremental Dataset Growing - DAgger

How can we make  $P_{\pi^*}(\mathbf{s}) \approx P_{\pi_\theta}(\mathbf{s})$ ?

**Key Idea** - If we cannot be clever on  $P_{\pi_\theta}(\mathbf{s})$  let us **be clever on  $P_{\pi^*}(\mathbf{s})$**

## DAgger – Dataset Aggregation

- 
1. Train  $\pi_\theta(a_t|s_t)$  from expert data  $\mathcal{D} = \{\mathbf{s}_1, \mathbf{a}_1, \dots, \mathbf{s}_N, \mathbf{a}_N\}$
  2. Run  $\pi_\theta(a_t|s_t)$  to get data  $\mathcal{D}_\pi = \{\mathbf{s}_1, \dots, \mathbf{s}_N\}$
  3. Ask expert to label  $\mathcal{D}_\pi$  with action  $\mathbf{a}_n$
  4. Aggregate  $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}_\pi$


### Potential issues

- Execute an unsafe/partially trained policy
- Repeatedly query the expert
- Expert is queried for an action without experiencing the state



# Dagger Demo


*Ross et al, Learning monocular reactive UAV control in cluttered natural environments, 2013*



<https://youtu.be/hNsP6-K3Hn4>

# Beyond Vanilla Dagger

- ✓ Experts do not need to be humans
  - ✓ Generative learning can be used for imitating expert policies
  - ✓ Solving **simpler optimization in a constrained part** of the state space
- ✓ Imitation then means **distilling knowledge of constrained policies into a general policy** that can do well in all scenarios



*Rusu, Policy distillation, ICLR 2016*

# Can we do better with expert data?

- Behavior cloning mimics the expert, but does not have a notion of **intention**
  - Expert may be suboptimal
  - Different embodiments
  - Robustness
- Effectively finding out what the teacher is trying to do could enable the agent to do even better than the demonstrator...

# Meeting the Expert Expectations

# Non-Markovian Behaviour

$$\pi_{\theta}(a_t | s_t)$$






If we see the same thing twice, we **do the same thing twice**, regardless of what happened before

$$\pi_{\theta}(a_t | s_t, s_{t-1}, \dots, s_0)$$




Behavior depends on the history of **past observations**

# Non-Markovian Behaviour – How to?




**Recurrent neural network**  
(gated to say the least)






# Multimodal Behaviour



- ✓ Avoiding an obstacle for an expert is easy
- ✓ Take one of the two steering angles

# Multimodal Behaviour in Regression

- ✓ Regression fails in multimodality
- ✓ MSE minimum is the average
- ✓ Which might not be advisable...



# Multimodal Behaviour in Regression



✓ ....unless you know how to do this

# Multimodality - Can we fix it?

## ✓ Autoregressive discretization

- ✓ Discretize the action space and train a classifier that predicts a categorical distribution it
- ✓ Do it progressively on the original features to avoid curse of dimensionality (e.g. PixelRNN)


## ✓ Gaussian mixture model output

- ✓ Predict mixture components weights, means and variances
- ✓ Need to guess number of modes


## ✓ Density model

- ✓ Density networks
- ✓ Variational Autoencoders
- ✓ Generative Adversarial Network

$$P(a_2|a_1)$$



$$P(a_1)$$



# Multimodality - Can we fix it?

## ✓ Autoregressive discretization

- ✓ Discretize the action space and train a classifier that predicts a categorical distribution it
- ✓ Do it progressively on the original features to avoid curse of dimensionality (e.g. PixelRNN)


## ✓ Gaussian mixture model output

- ✓ Predict mixture components weights, means and variances
- ✓ Need to guess number of modes

## ✓ Density model

- ✓ Density networks
- ✓ Variational Autoencoders
- ✓ Generative Adversarial Network

$$\pi_{\theta}(s|a) = \sum_i w_i \mathcal{N}(\mu_i, \Sigma_i)$$



# Multimodality - Can we fix it?

## ✓ Autoregressive discretization


- ✓ Discretize the action space and train a classifier that predicts a categorical distribution it
- ✓ Do it progressively on the original features to avoid curse of dimensionality (e.g. PixelRNN)

## ✓ Gaussian mixture model output

- ✓ Predict mixture components weights, means and variances
- ✓ Need to guess number of modes

## ✓ Density model


- ✓ Density networks
- ✓ Variational Autoencoders
- ✓ Generative Adversarial Network



# Generative Imitation Learning



# Generative Adversarial Imitation Learning (GAIL)



- ✓ Use a policy network as generator (state conditioned)
- ✓ Find a policy that makes it impossible for a discriminator network to distinguish between state-actions from the expert demonstrations and state-actions visited by the learnt policy
- ✓ Generator needs to be trained using RL

Ho and Ermon, Generative Adversarial Imitation Learning, NIPS 2016

# GAIL Algorithm

---

## Algorithm 1 Generative adversarial imitation learning

---

- 1: **Input:** Expert trajectories  $\tau_E \sim \pi_E$ , initial policy and discriminator parameters  $\theta_0, w_0$
- 2: **for**  $i = 0, 1, 2, \dots$  **do**
- 3:   Sample trajectories  $\tau_i \sim \pi_{\theta_i}$
- 4:   Update the discriminator parameters from  $w_i$  to  $w_{i+1}$  with the gradient

$$\hat{\mathbb{E}}_{\tau_i} [\nabla_w \log(D_w(s, a))] + \hat{\mathbb{E}}_{\tau_E} [\nabla_w \log(1 - D_w(s, a))] \quad (17)$$

- 5:   Take a policy step from  $\theta_i$  to  $\theta_{i+1}$ , using the TRPO rule with cost function  $\log(D_{w_{i+1}}(s, a))$ . Specifically, take a KL-constrained natural gradient step with

$$\hat{\mathbb{E}}_{\tau_i} [\nabla_{\theta} \log \pi_{\theta}(a|s) Q(s, a)] - \lambda \nabla_{\theta} H(\pi_{\theta}), \quad (18)$$

where  $Q(\bar{s}, \bar{a}) = \hat{\mathbb{E}}_{\tau_i} [\log(D_{w_{i+1}}(s, a)) \mid s_0 = \bar{s}, a_0 = \bar{a}]$

- 6: **end for**
- 

Discriminator  
provides reward  
function

Entropy-based  
policy regularizer

Trust-Region policy gradient

# Rewards are not always as explicit

reward



Mnih et al. '15


robotics



dialog



autonomous driving



what is the **reward**?  
often use a proxy

# Inverse Reinforcement Learning (IRL)

- ✓ An alternative to imitation learning
  - ✓ Use **demonstrations to learn a reward** function
  - ✓ Train a **policy using learnt reward** function
- ✓ Least expensive form of supervision
  - ✓ Does **not need full demonstrations**
  - ✓ RL phase can “fill in” missing behavior given partial demonstrations
- ✓ Argued to be a more comprehensive model of expert behavior
  - ✓ Learning **why the expert did something instead of mapping states to actions**
  - ✓ Can potentially generalize better

GAIL is a particular form of inverse RL that learns a reward function that tries to match state distributions between the expert and imitator

# More Formally


## Forward Reinforcement Learning

- ✓ Given
  - ✓ States  $s \in \mathcal{S}$  and actions  $a \in \mathcal{A}$
  - ✓ Transitions  $P_{ss'}^a$ , (sometimes)
  - ✓ Reward function  $\mathcal{R}_s^a$
- ✓ Learn or infer policy  $\pi^*(a|s)$

## Inverse Reinforcement Learning

- ✓ Given
  - ✓ States  $s \in \mathcal{S}$  and actions  $a \in \mathcal{A}$
  - ✓ Transitions  $P_{ss'}^a$ , (sometimes)
  - ✓ Sampled episodes from expert  
 $(s, a) \sim \pi^*(a|s)$
- ✓ Learn a reward function  $r_\phi(a, s)$ , with  $\phi$  being adaptive parameters
- ✓ ...and use  $r_\phi(a, s)$  to learn/infer  $\pi^*(a|s)$

# Solving IRL – MaxEntropy Deep IRL



## Algorithm 1 Maximum Entropy Deep IRL

**Input:**  $\mu_D^a, f, S, A, T, \gamma$

**Output:** optimal weights  $\theta^*$  ← Expert state-action frequencies

1:  $\theta^1 = \text{initialise\_weights}()$

### Iterative model refinement

2: **for**  $n = 1 : N$  **do**

3:  $r^n = \text{nn\_forward}(f, \theta^n)$

### Solution of MDP with current reward

4:  $\pi^n = \text{approx\_value\_iteration}(r^n, S, A, T, \gamma)$

5:  $\mathbb{E}[\mu^n] = \text{propagate\_policy}(\pi^n, S, A, T)$

### Determine Maximum Entropy loss and gradients

6:  $\mathcal{L}_D^n = \log(\pi^n) \times \mu_D^a$

7:  $\frac{\partial \mathcal{L}_D^n}{\partial r^n} = \mu_D - \mathbb{E}[\mu^n]$

### Compute network gradients

8:  $\frac{\partial \mathcal{L}_D^n}{\partial \theta^n} = \text{nn\_backprop}(f, \theta^n, \frac{\partial \mathcal{L}_D^n}{\partial r^n})$

9:  $\theta^{n+1} = \text{update\_weights}(\theta^n, \frac{\partial \mathcal{L}_D^n}{\partial \theta^n})$

10: **end for**

Finds  $Q$  and  $V$  and infers  $\pi^n$


Expected state visiting frequencies by sampling with  $\pi^n$

MAP of observing expert samples

Wulfmeier et al, Maximum Entropy Deep Inverse Reinforcement Learning, 2015



# Guided Cost Learning



Finn et al, Guided Cost Learning: Deep Inverse Optimal Control via Policy Optimization, ICML 2016



# GCL Demo

## Guided Cost Learning: Deep Inverse Optimal Control via Policy Optimization

Chelsea Finn, Sergey Levine, Pieter Abbeel  
UC Berkeley

<https://youtu.be/hXxaepw0zAw>

# Take home messages

- ✓ Effective **imitation learning** is much about managing distribution shift and relaying less on human demonstration
- ✓ Using a **learnable model of reward** can serve the purpose of reducing the extent of human labelling (**inverse reinforcement learning**)
- ✓ Much left unsaid
  - ✓ Off-policy learning from imitation policy
  - ✓ Q-learning as natural off-policy imitation learner
    - ✓ Just drop demonstrations into the replay buffer
  - ✓ Inverse reinforcement learning Vs generative adversarial learning